

Task 1: Understanding Process Creation and Scheduling

Objective: Learn how processes are created and scheduled in an operating system.

Steps:

1. Create a simple "Hello World" program in your preferred programming language (e.g., C, Python, Java).
2. Compile and run the program, observing the process ID (PID) and other information using system tools like `ps` or `top` in Linux.
3. Investigate how the OS schedules your process for execution by checking its state (e.g., ready, running) during execution using `ps` or other process monitoring tools.
4. Create multiple processes and observe how they are handled by the OS.

Question: How does the operating system decide which process gets CPU time? What role does the process scheduler play in this?

Task 2: Exploring Thread Management and Concurrency

Objective: Understand the concept of multithreading and how concurrency works in an operating system.

Steps:

1. Write a simple program that uses multiple threads to perform tasks concurrently (e.g., one thread for downloading a file, another for processing data).
2. Use thread management tools to monitor and manage these threads (e.g., `top`, `htop`, or task manager in your OS).
3. Introduce synchronization mechanisms like semaphores or mutexes to manage shared resources between threads.
4. Test how threads interact with each other and ensure data consistency is maintained with synchronization.

Question: How does multithreading improve system performance, and what challenges does it introduce in terms of data consistency and synchronization?

Task 3: Investigating Deadlock in Multithreading

Objective: Learn how deadlock occurs in multithreading environments and how to detect and resolve it.

Steps:

1. Write a simple program with two threads where each thread locks a resource, then tries to acquire the lock on the other's resource, causing a deadlock.
2. Use a debugger or logging to trace the execution and confirm that the program is deadlocked.

3. Modify the program to resolve the deadlock by introducing timeouts or changing the locking order.

Question: What conditions are necessary for a deadlock to occur? How can deadlock be avoided or resolved in multithreaded programs?

Task 4: Understanding I/O Operations and Buffering

Objective: Explore how the OS handles I/O operations using buffers and how data is managed between processes and devices.

Steps:

1. Write a program that reads from and writes to a file, ensuring that the data is buffered for efficiency.
2. Use system tools like `vmstat` or `iostat` to monitor I/O operations and buffer usage while your program is running.
3. Introduce large files to your program and observe how the OS handles large-scale I/O operations, paying attention to any performance impacts.

Question: How do I/O buffers improve the efficiency of file operations? What impact do buffer sizes have on the system's performance during I/O operations?

Task 5: Analyzing I/O Scheduling Algorithms

Objective: Understand how different I/O scheduling algorithms impact system performance.

Steps:

1. Use a Linux-based system and run multiple I/O intensive tasks like copying large files or running disk benchmarks.
2. Analyze I/O scheduling using the `iostat` or `dstat` command to monitor disk activity and the types of scheduling algorithms in use (e.g., CFQ, Deadline, or NOOP).
3. Test different algorithms and measure the time taken for I/O operations, then compare the results.

Question: How do different I/O scheduling algorithms affect disk performance? What are the advantages and disadvantages of each algorithm in different workloads?

Task 6: Exploring the OSI Model - Layers and Functions

Objective: Learn about the OSI model and the role each layer plays in network communication.

Steps:

1. Research and create a diagram of the OSI model, labeling each of the seven layers and describing their functions.
 2. Set up a basic network communication system using a protocol like **ping** (Layer 3) or **HTTP** (Layer 7).
 3. Use network tools like `tracert` or `tcpdump` to capture packets and identify which OSI layers are involved during communication.
- Question:** How does the OSI model help in understanding network communication? What role does each layer play in ensuring data is transmitted across the network?
-

Task 7: Testing Transport Layer Protocols (TCP/UDP)

Objective: Understand the differences between TCP and UDP and how they function in the transport layer.

Steps:

1. Use the `netstat` command to observe network connections and identify which ones use TCP and which use UDP.
 2. Write simple client-server programs using TCP and UDP (e.g., a chat application).
 3. Test both programs and observe how TCP guarantees reliability (via handshakes and acknowledgments) while UDP does not.
- Question:** What are the key differences between TCP and UDP in terms of reliability and performance? When would you choose one over the other?
-

Task 8: Exploring Linux Process Management

Objective: Understand how processes are managed in Linux, including creation, scheduling, and termination.

Steps:

1. Use the `ps` and `top` commands to list the running processes on a Linux system.
 2. Create and kill processes using the `fork()` system call or equivalent in a C program.
 3. Monitor process states using `ps` or `htop` to understand how the Linux kernel schedules processes.
- Question:** How does Linux manage processes in terms of scheduling, and how can you observe and control the state of processes using system tools?
-

Task 9: Linux File System and Permissions

Objective: Learn how Linux organizes files and manages permissions for users.

Steps:

1. Explore the Linux file system using commands like `ls`, `cd`, and `df` to understand the directory structure and disk usage.
2. Create files and directories, setting various permissions using the `chmod`, `chown`, and `chgrp` commands.
3. Test access control by creating multiple users and groups and modifying file permissions to restrict or allow access.

Question: How do file permissions in Linux enhance security? How does the ownership model (user, group, others) control access to resources?

Task 10: Using Shell Scripting for Process Automation

Objective: Learn how to automate system tasks using Linux shell scripting.
Steps:

1. Write a shell script to automate a simple task, such as backing up a directory or renaming files in a batch process.
2. Use system commands like `cron` to schedule the script to run at a specific time or interval.
3. Test the script and ensure it runs as expected, handling errors or exceptions if necessary.

Question: How can shell scripting help automate system management tasks? What advantages does it offer for managing processes and resources on a Linux system?